

This notebook contains codes executed for Assignment-4 (**Predictive Data Mining**) for the course **CSIT558_01SP25** by:

- Chaudhari, Amishaben Natavarbhai (CWID:50129394)
- Javed, Syed Mehrose (CWID: 50128848)
- Subhani, Muhammad Kamal (CWID: 50136417)
- Thangavel Sathiyamoorthy, Punithan (CWID: 50135877)
- Venkata Appala Manoj Muvvala (CWID: 50122981)

```
# Importing the required libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeRegressor
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import mean_absolute_error, mean_squared_error,
r2_score
import joblib

import warnings
warnings.filterwarnings("ignore")
```

DATA LOADING AND EDA

```
# Loading the data
data = pd.read_csv("drug1000.csv")
data

{"summary": {"name": "data", "rows": 1000, "fields": [{"column": "Age", "properties": {"dtype": "number", "std": 16, "min": 15, "max": 79, "num_unique_values": 65, "samples": [62, 73, 29]}, "semantic_type": "\\", "description": "\\n"}, {"column": "Sex", "properties": {"dtype": "category", "num_unique_values": 2, "samples": ["M", "F"], "semantic_type": "\\", "description": "\\n"}, {"column": "BP", "properties": {"dtype": "category", "num_unique_values": 3, "samples": ["LOW", "NORMAL", "HIGH"], "semantic_type": "\\", "description": "\\n"}, {"column": "Cholesterol", "properties": {"dtype": "category", "num_unique_values": 2, "samples": ["NORMAL", "HIGH"], "semantic_type": "\\", "description": "\\n"}, {"column": "Na_to_K", "properties": {"dtype": "category", "samples": ["\\n"], "semantic_type": "\\", "description": "\\n"}], "semantic_type": "\\", "description": "\\n"}]}
```

```

{"\n      "dtype": "number",\n      "std":\n7.652179968263982,\n      "min": 5.921618772385234,\n      "max": 41.33547607149792,\n      "num_unique_values": 1000,\n      "samples": [\n          10.38522916646769,\n29.611525616870747\n      ],\n      "semantic_type": "\\",\\n\n      "description": \"\\n      \"},\\n      {\n          "column":\n          "Drug",\\n\n          "properties": {\n              "dtype": "category",\\n\n              "num_unique_values": 5,\n              "samples": [\n                  "drugY",\\n\n                  "drugA"\n              ],\n              "semantic_type": "\",\\n\n              "description": \"\\n      \"}\\n\n      }\n  ],\n  "type": "dataframe",\n  "variable_name": "data"
}

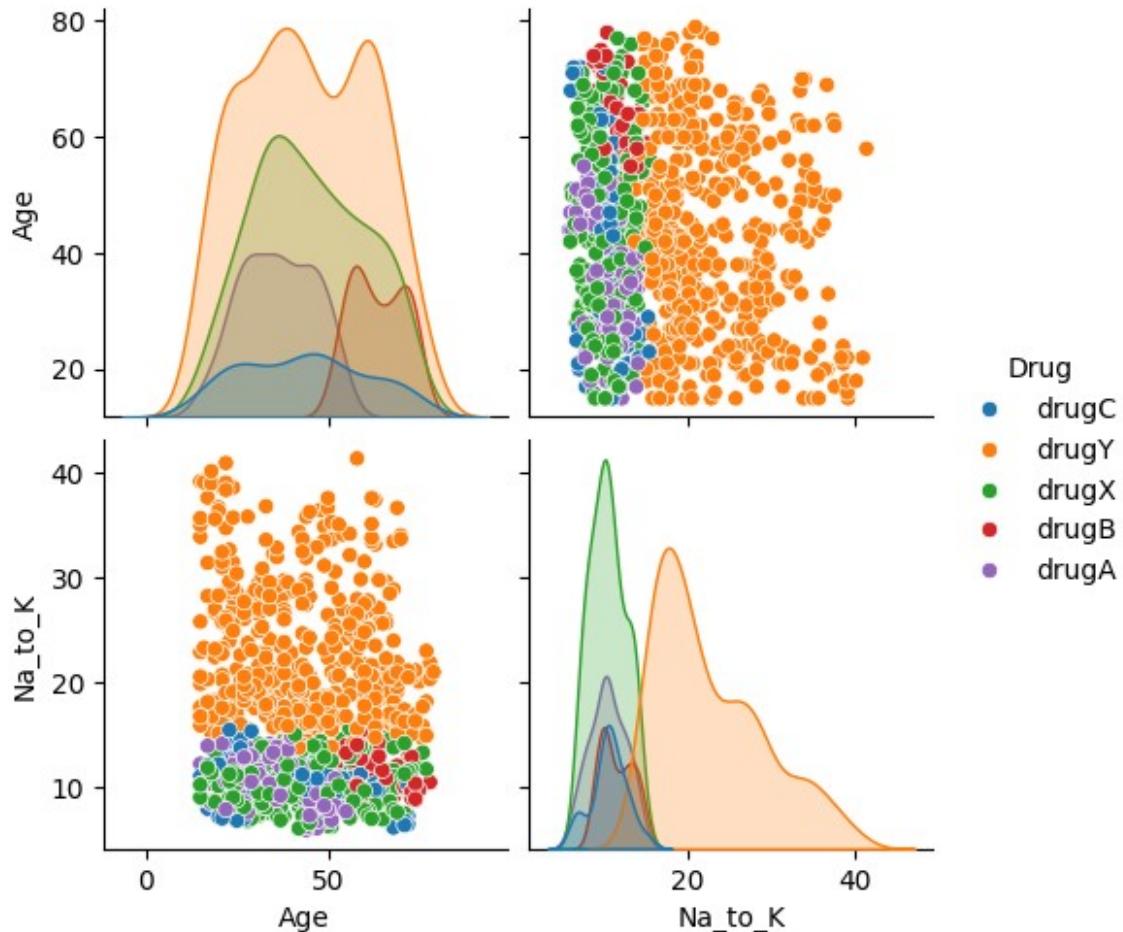
# Exploring the data
data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 6 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   Age         1000 non-null    int64  
 1   Sex         1000 non-null    object  
 2   BP          1000 non-null    object  
 3   Cholesterol 1000 non-null   object  
 4   Na_to_K     1000 non-null    float64 
 5   Drug        1000 non-null    object  
dtypes: float64(1), int64(1), object(4)
memory usage: 47.0+ KB

# Plotting a pair-plot of the attributes in the dataset
sns.pairplot(data=data, hue='Drug')

<seaborn.axisgrid.PairGrid at 0x79b6221add90>

```



DATA CLEANING

```
# Checking the data for null values
data.isnull().sum()

Age          0
Sex          0
BP           0
Cholesterol 0
Na_to_K      0
Drug         0
dtype: int64

# Checking for duplicates
duplicates = data.duplicated().sum()
print(f"Number of duplicate rows: {duplicates}")

Number of duplicate rows: 0

# checking the data by printing the initial 5 rows
data.head()
```

```

{
  "summary": {
    "name": "data",
    "rows": 1000,
    "fields": [
      {
        "column": "Age",
        "properties": {
          "dtype": "number",
          "std": 16,
          "min": 15,
          "max": 79,
          "num_unique_values": 65,
          "samples": [62, 73, 29]
        },
        "semantic_type": "\",
        "description": "\n"
      },
      {
        "column": "Sex",
        "properties": {
          "dtype": "category",
          "num_unique_values": 2,
          "samples": ["M", "F"]
        },
        "semantic_type": "\",
        "description": "\n"
      },
      {
        "column": "BP",
        "properties": {
          "dtype": "category",
          "num_unique_values": 3,
          "samples": ["LOW", "NORMAL", "HIGH"]
        },
        "semantic_type": "\",
        "description": "\n"
      },
      {
        "column": "Cholesterol",
        "properties": {
          "dtype": "category",
          "num_unique_values": 2,
          "samples": ["NORMAL", "HIGH"]
        },
        "semantic_type": "\",
        "description": "\n"
      },
      {
        "column": "Na_to_K",
        "properties": {
          "dtype": "number",
          "std": 7.652179968263982,
          "min": 5.921618772385234,
          "max": 41.33547607149792,
          "num_unique_values": 1000,
          "samples": [10.38522916646769, 29.611525616870747]
        },
        "semantic_type": "\",
        "description": "\n"
      }
    ],
    "column": "Drug",
    "properties": {
      "dtype": "category",
      "num_unique_values": 5,
      "samples": ["drugC", "drugY", "drugX", "drugB", "drugA"]
    },
    "semantic_type": "\",
    "description": "\n"
  }
}

```

MODELING PREPARATION

```

#Printing the unique values of the categorical variables
print(data['Sex'].unique())
print(data['BP'].unique())
print(data['Cholesterol'].unique())
print(data['Drug'].unique())

['F' 'M']
['LOW' 'NORMAL' 'HIGH']
['HIGH' 'NORMAL']
['drugC' 'drugY' 'drugX' 'drugB' 'drugA']

# Mapping the categorical variables to numerical values
data['Sex'] = data['Sex'].replace({'F':1, 'M':2})
data['BP'] = data['BP'].replace({'LOW':1, 'NORMAL':2, 'HIGH':3})
data['Cholesterol'] = data['Cholesterol'].replace({'LOW':1, 'NORMAL':2, 'HIGH':3})

```

```

data['Drug'] = data['Drug'].replace({'drugA':1, 'drugB':2, 'drugC':3,
'drugX':4, 'drugY':5})
data.head()

{
  "summary": {
    "name": "data",
    "rows": 1000,
    "fields": [
      {
        "column": "Age",
        "properties": {
          "dtype": "number",
          "std": 16,
          "min": 15,
          "max": 79,
          "num_unique_values": 65,
          "samples": [62, 73, 29]
        }
      },
      {
        "column": "Sex",
        "properties": {
          "dtype": "number",
          "std": 0,
          "min": 1,
          "max": 2,
          "num_unique_values": 2,
          "samples": [2, 1]
        }
      },
      {
        "column": "BP",
        "properties": {
          "dtype": "number",
          "std": 0,
          "min": 1,
          "max": 3,
          "num_unique_values": 3,
          "samples": [1, 2]
        }
      },
      {
        "column": "Cholesterol",
        "properties": {
          "dtype": "number",
          "std": 0,
          "min": 2,
          "max": 3,
          "num_unique_values": 2,
          "samples": [2, 3]
        }
      },
      {
        "column": "Na_to_K",
        "properties": {
          "dtype": "number",
          "std": 7.652179968263982,
          "min": 5.921618772385234,
          "max": 41.33547607149792,
          "num_unique_values": 1000,
          "samples": [10.38522916646769, 29.611525616870747]
        }
      }
    ]
  }
}

# Separating the target variable from the rest of the variables
features = ["Age", "Sex", "BP", "Cholesterol", "Na_to_K"]

x = data[features]
y = data.Drug

# Checking the variables other than the target variable
x.head()

{
  "summary": {
    "name": "x",
    "rows": 1000,
    "fields": [
      {
        "column": "Age",
        "properties": {
          "dtype": "number",
          "std": 16,
          "min": 15,
          "max": 79
        }
      }
    ]
  }
}

```

```

    \\"max\": 79,\n          \\"num_unique_values\": 65,\n          \\"samples\": [\n            62,\n            73,\n            29\n          ],\n        },\n        \\"semantic_type\": \"\",\n        \\"description\": \"\",\n        \\"column\": \"Sex\",\n        \\"properties\": {\n          \\"dtype\\\": \"number\",\n          \\"std\\\": 0,\n          \\"min\\\": 1,\n          \\"samples\\\": 2,\n          \\"num_unique_values\\\": 2,\n          \\"semantic_type\\\": \"\",\n          \\"description\\\": \"\",\n          \\"column\\\": \"BP\",\n          \\"properties\\\": {\n            \\"number\\\", \n            \\"std\\\": 0,\n            \\"min\\\": 1,\n            \\"max\\\": 3,\n            \\"num_unique_values\\\": 3,\n            \\"samples\\\": [\n              1,\n              2\n            ],\n            \\"semantic_type\\\": \"\",\n            \\"description\\\": \"\",\n            \\"column\\\": \"Cholesterol\",\n            \\"properties\\\": {\n              \\"dtype\\\": \"number\",\n              \\"std\\\": 0,\n              \\"min\\\": 2,\n              \\"max\\\": 3,\n              \\"num_unique_values\\\": 2,\n              \\"samples\\\": [\n                2,\n                3\n              ],\n              \\"semantic_type\\\": \"\",\n              \\"description\\\": \"\",\n              \\"column\\\": \"Na_to_K\",\n              \\"properties\\\": {\n                \\"number\\\", \n                \\"std\\\": 7.652179968263982,\n                \\"min\\\": 5.921618772385234,\n                \\"max\\\": 41.33547607149792,\n                \\"num_unique_values\\\": 1000,\n                \\"samples\\\": [\n                  10.38522916646769,\n                  29.611525616870747\n                ],\n                \\"semantic_type\\\": \"\",\n                \\"description\\\": \"\"\n              }\n            },\n            \"type\": \"dataframe\", \"variable_name\": \"x\"}\n\n# Checking the target variable\nv.head()

```

```
0    3  
1    5  
2    5  
3    4  
4    4  
Name: Drug, dtype: int64
```

```
# Splitting the data into Train and Test Sets  
train_x, test_x, train_y, test_y = train_test_split(x, y, random_state  
= 20)
```

TUNING THE HYPERPARAMETERS

Tuning Hyperparameters: max depth

Training models with different max depth values

depths = [4, 5, 6, 7, 8, 10, 20]

```
models = {depth: DecisionTreeClassifier(max_depth=depth,
random_state=20).fit(train_x, train_y) for depth in depths}
```

Making predictions

```

predictions = {depth: models[depth].predict(test_x) for depth in depths}

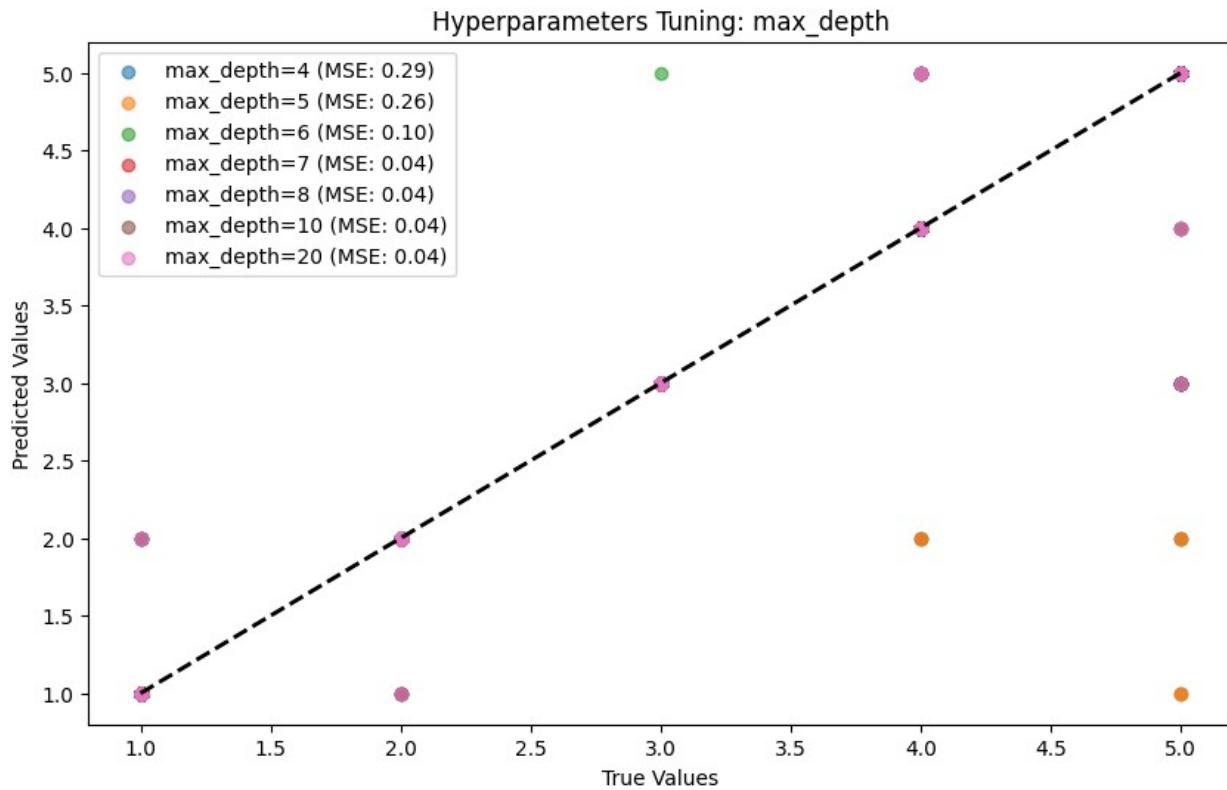
# Evaluating performance using mean_squared_error
mse_scores = {depth: mean_squared_error(test_y, predictions[depth])
for depth in depths}

# Visualizing True vs Predicted values
plt.figure(figsize=(10, 6))

for depth in depths:
    # Plotting the MSE's for different max_depth values
    plt.scatter(test_y, predictions[depth], alpha=0.6,
label=f"max_depth={depth} (MSE: {mse_scores[depth]:.2f})")

plt.plot([min(test_y), max(test_y)], [min(test_y), max(test_y)],
'k--', lw=2) # Perfect prediction line
plt.xlabel("True Values")
plt.ylabel("Predicted Values")
plt.title("Hyperparameters Tuning: max_depth")
plt.legend()
plt.show()

```



```

# Tuning Hyperparameters: max_leaf_nodes

```

```

# Training models with different max_leaf_nodes values
leafnodes = [5, 8, 10, 20, 25, 30, 50, 75]
models = {leafnode: DecisionTreeClassifier(max_leaf_nodes=leafnode,
random_state=20).fit(train_x, train_y) for leafnode in leafnodes}

# Making predictions
predictions = {leafnode: models[leafnode].predict(test_x) for leafnode in leafnodes}

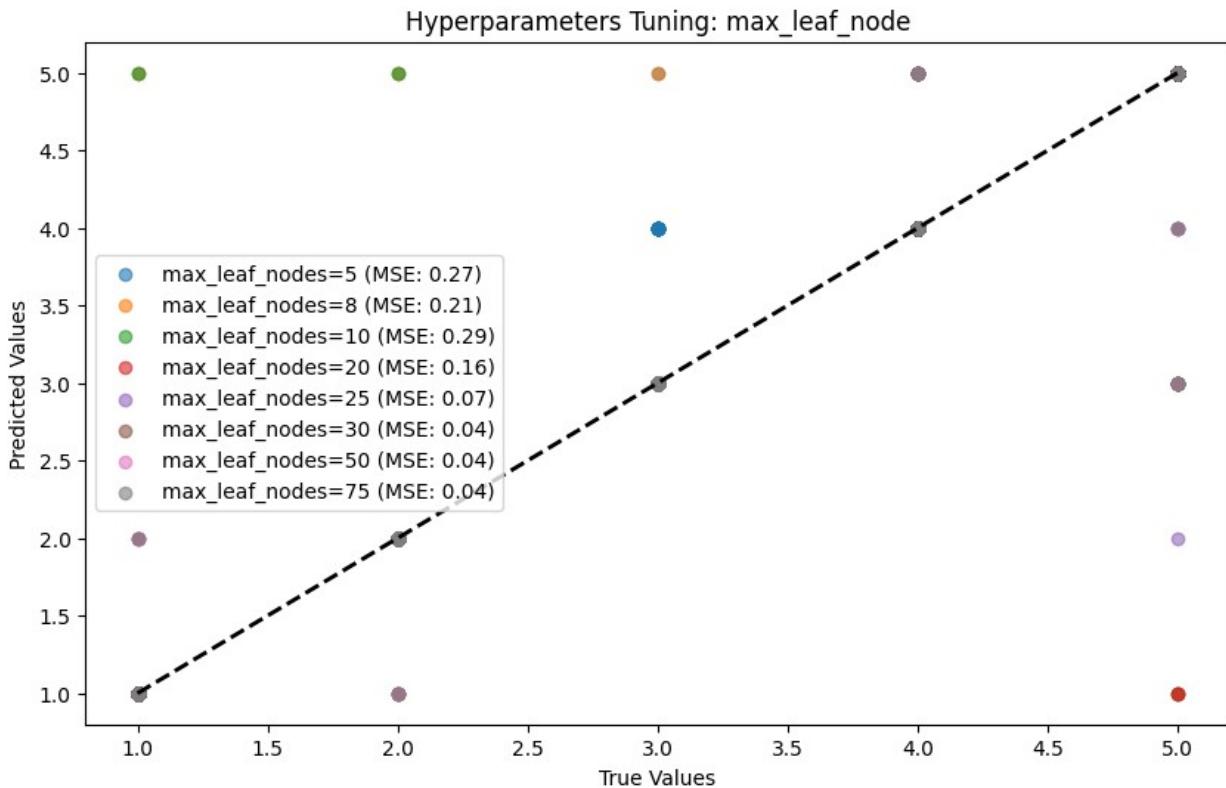
# Evaluating performance using mean_squared_error
mse_scores = {leafnode: mean_squared_error(test_y,
predictions[leafnode]) for leafnode in leafnodes}

# Visualizing True vs Predicted values
plt.figure(figsize=(10, 6))

for leafnode in leafnodes:
    # Plotting the MSE's for different leaf nodes
    plt.scatter(test_y, predictions[leafnode], alpha=0.6,
label=f"max_leaf_nodes={leafnode} (MSE: {mse_scores[leafnode]:.2f})")

plt.plot([min(test_y), max(test_y)], [min(test_y), max(test_y)],
'k--', lw=2) # Perfect prediction line
plt.xlabel("True Values")
plt.ylabel("Predicted Values")
plt.title("Hyperparameters Tuning: max_leaf_node")
plt.legend()
plt.show()

```



```
# Tuning Hyperparameters : max_features

# Training models with different max_leaf_nodes values
feature_settings = [None, 'sqrt', 'log2']
models = {feature: DecisionTreeClassifier(max_features=feature,
random_state=20).fit(train_x, train_y) for feature in
feature_settings}

# Making predictions
predictions = {feature: models[feature].predict(test_x) for feature in
feature_settings}

# Evaluating performance using mean_squared_error
mse_scores = {feature: mean_squared_error(test_y,
predictions[feature]) for feature in feature_settings}

# Visualizing True vs Predicted values
plt.figure(figsize=(10, 6))

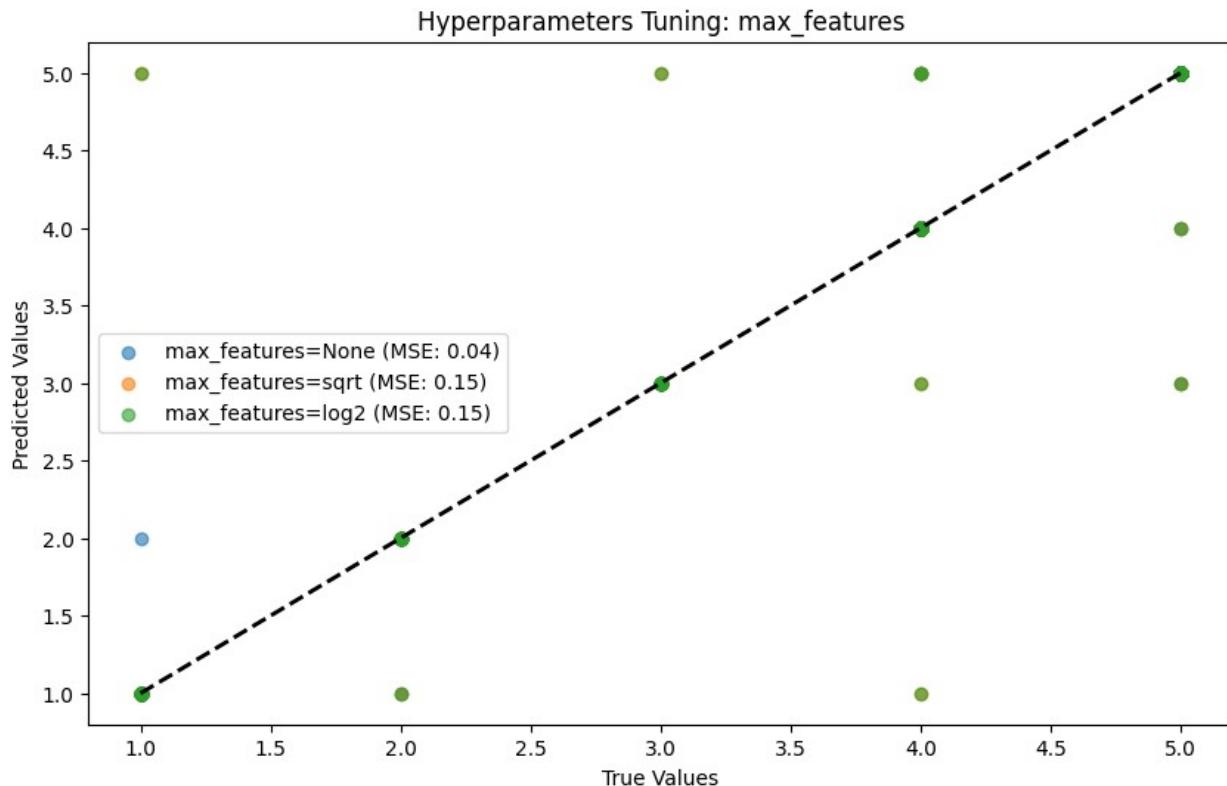
for feature in feature_settings:
    plt.scatter(test_y, predictions[feature], alpha=0.6,
label=f"max_features={feature} (MSE: {mse_scores[feature]:.2f})") # Plotting the MSE's for different leaf nodes

plt.plot([min(test_y), max(test_y)], [min(test_y), max(test_y)],
'k--', lw=2) # Perfect prediction line
```

```

plt.xlabel("True Values")
plt.ylabel("Predicted Values")
plt.title("Hyperparameters Tuning: max_features")
plt.legend()
plt.show()

```



CREATING THE FINAL CLASSIFIER MODEL BASED ON THE MOST OPTIMAL VALUES OF THE HYPERPARAMETERS

```

# Training and fitting the Decision Tree Classifier model based on the tuned hyperparameters
model = DecisionTreeClassifier(max_depth=7, max_leaf_nodes = 30,
max_features=None, random_state = 20)
model.fit(train_x, train_y)

DecisionTreeClassifier(max_depth=7, max_leaf_nodes=30,
random_state=20)

# Getting model parameters
model.get_params()

{'ccp_alpha': 0.0,
 'class_weight': None,
 'criterion': 'gini',
 'max_depth': 7,
 'max_features': None,

```

```

'max_leaf_nodes': 30,
'min_impurity_decrease': 0.0,
'min_samples_leaf': 1,
'min_samples_split': 2,
'min_weight_fraction_leaf': 0.0,
'monotonic_cst': None,
'random_state': 20,
'splitter': 'best'}

# Verifying the model by checking the error
y_pred = model.predict(test_x)

mae = mean_absolute_error(test_y, y_pred)
mse = mean_squared_error(test_y, y_pred)
r2 = r2_score(test_y, y_pred)

print(f'MAE: {mae:.2f}')
print(f'MSE: {mse:.2f}')
print(f'R2: {r2:.2f}')

MAE: 0.03
MSE: 0.04
R2: 0.98

```

EXPERIMENTING WITH OTHER ACCEPTABLE VALUES OF THE HYPERPARAMETERS

WHEN max_depth=7, max_leaf_nodes = 25, max_features='log2'

```

# Training and fitting the model when
# max_depth=7, max_leaf_nodes = 25, max_features='log2'
model = DecisionTreeClassifier(max_depth=7,
                               max_leaf_nodes = 25,
                               max_features='log2', random_state = 20)
model.fit(train_x, train_y)

# Verifying the model by checking the error
y_pred = model.predict(test_x)

mae = mean_absolute_error(test_y, y_pred)
mse = mean_squared_error(test_y, y_pred)
r2 = r2_score(test_y, y_pred)

print(f'MAE: {mae:.2f}')
print(f'MSE: {mse:.2f}')
print(f'R2: {r2:.2f}')

MAE: 0.12
MSE: 0.25
R2: 0.87

```

```

max_depth=10, max_leaf_nodes = 50, max_features='sqrt'

# Traiing and fitting the model when
# max_depth=10, max_leaf_nodes = 50, max_features='sqrt'
model = DecisionTreeClassifier(max_depth=10,
                               max_leaf_nodes = 50,
                               max_features='sqrt', random_state = 20)
model.fit(train_x, train_y)

# Verifying the model by checking the error
y_pred = model.predict(test_x)

mae = mean_absolute_error(test_y, y_pred)
mse = mean_squared_error(test_y, y_pred)
r2 = r2_score(test_y, y_pred)

print(f'MAE: {mae:.2f}')
print(f'MSE: {mse:.2f}')
print(f'R2: {r2:.2f}')

MAE: 0.11
MSE: 0.28
R2: 0.86

```

WHEN max_depth=8, max_leaf_nodes = 20, max_features='log2'

```

# Traiing and fitting the model when
# max_depth=8, max_leaf_nodes = 20, max_features='log2'
model = DecisionTreeClassifier(max_depth=8,
                               max_leaf_nodes = 20,
                               max_features='log2', random_state = 20)
model.fit(train_x, train_y)

# Verifying the model by checking the error
y_pred = model.predict(test_x)

mae = mean_absolute_error(test_y, y_pred)
mse = mean_squared_error(test_y, y_pred)
r2 = r2_score(test_y, y_pred)

print(f'MAE: {mae:.2f}')
print(f'MSE: {mse:.2f}')
print(f'R2: {r2:.2f}')

MAE: 0.21
MSE: 0.47
R2: 0.76

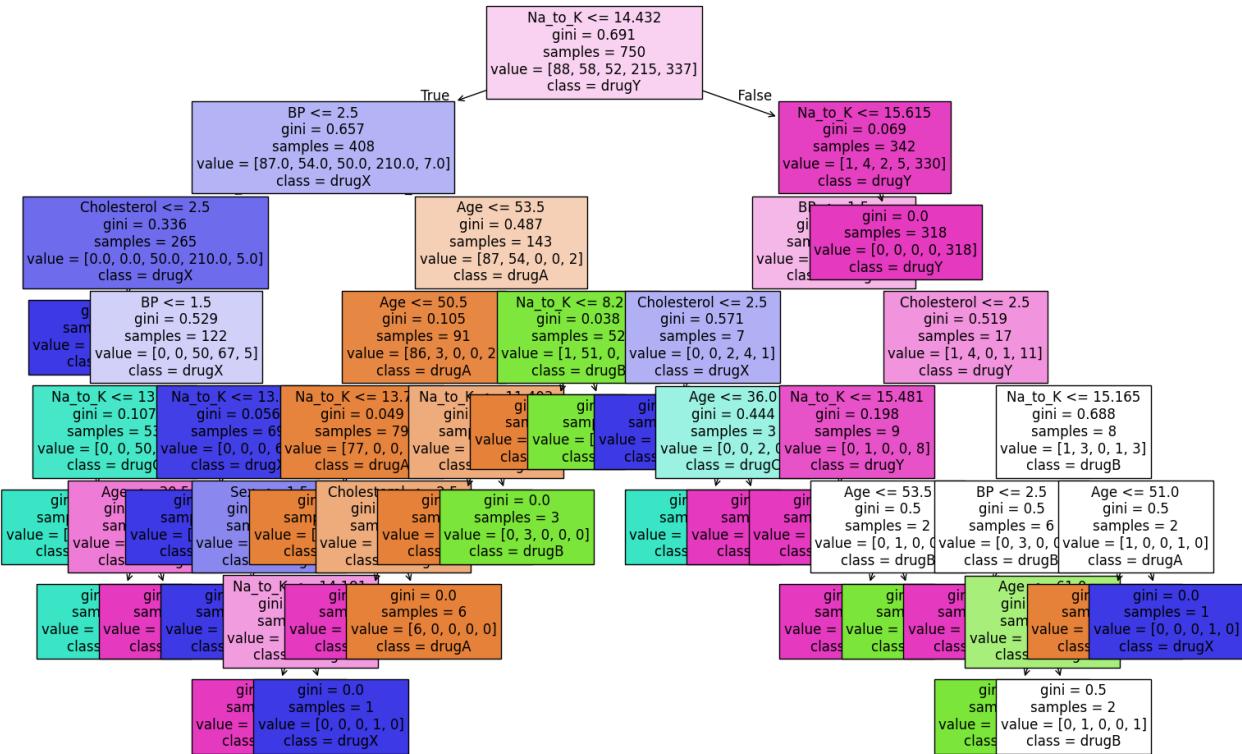
```

PLOTTING THE DECISION TREE

```
# Plotting the tree in diagramtic representation
from sklearn import tree

fig = plt.figure(figsize=(18, 12))
_ = tree.plot_tree(model,
                    feature_names=x.columns.tolist(),
                    class_names=['drugA', 'drugB', 'drugC', 'drugX',
                    'drugY'],
                    filled=True,
                    fontsize=12
)

```



CREATING THE CLASSIFIER TOOL

```
# Storing the trained model as pickle file
joblib.dump(model, 'drug_classifier.pkl')

['drug_classifier.pkl']

# Loading the trained model
model = joblib.load('drug_classifier.pkl')

# Defining the mappings for categorical variables
sex_map = {'M': 1, 'F': 2}
```

```

bp_map = {'LOW': 1, 'NORMAL': 2, 'HIGH': 3}
cholesterol_map = {'LOW': 1, 'NORMAL': 2, 'HIGH': 3}

# Defining the function to predict the drug
def predict_drug(age, sex, bp, cholesterol, na_to_k):
    # Converting categorical inputs to numerical values
    sex = sex_map.get(sex.upper(), 0)
    bp = bp_map.get(bp.upper(), 0)
    cholesterol = cholesterol_map.get(cholesterol.upper(), 0)

    # Preparing input array
    features = np.array([[age, sex, bp, cholesterol, na_to_k]])

    # Making prediction
    prediction = model.predict(features)[0]

    # Mapping the Predicted numerical value to its corresponding Drug Name
    drug_map = {1: 'drugA', 2: 'drugB', 3: 'drugC', 4: 'drugX', 5: 'drugY'}
    prediction = drug_map.get(prediction, 'Unknown')

    return prediction

```

Testing the Prediction with different sets of inputs

```

# Input: 56 F LOW HIGH 11.567 drugC
age = 56
sex = "F"
bp = "LOW"
cholesterol = "HIGH"
na_to_k = 11.567

# Getting the Prediction
predicted_drug = predict_drug(age, sex, bp, cholesterol, na_to_k)

# Printing the Result
print(f"Predicted Drug: {predicted_drug}")

Predicted Drug: drugC

# Input: 85 M HIGH LOW 46.5 drugY
age = 85
sex = "M"
bp = "HIGH"
cholesterol = "LOW"
na_to_k = 46.5

# Getting the Prediction
predicted_drug = predict_drug(age, sex, bp, cholesterol, na_to_k)

```

```
# Printing the Result
print(f"Predicted Drug: {predicted_drug}")

Predicted Drug: drugY

# Input: 75 M NORMAL NORMAL 12.33 drugX
age = 75
sex = "M"
bp = "NORMAL"
cholesterol = "NORMAL"
na_to_k = 12.33

# Getting the Prediction
predicted_drug = predict_drug(age, sex, bp, cholesterol, na_to_k)

# Printing the Result
print(f"Predicted Drug: {predicted_drug}")

Predicted Drug: drugX

# Input: 19 F HIGH HIGH 13.313 drugA
age = 19
sex = "F"
bp = "HIGH"
cholesterol = "HIGH"
na_to_k = 13.313

# Getting the Prediction
predicted_drug = predict_drug(age, sex, bp, cholesterol, na_to_k)

# Printing the Result
print(f"Predicted Drug: {predicted_drug}")

Predicted Drug: drugA

# Input: 60 F HIGH HIGH 13.303 drugB
age = 60
sex = "F"
bp = "HIGH"
cholesterol = "HIGH"
na_to_k = 13.303

# Getting the Prediction
predicted_drug = predict_drug(age, sex, bp, cholesterol, na_to_k)

# Printing the Result
print(f"Predicted Drug: {predicted_drug}")

Predicted Drug: drugB
```

INTERACTIVE INPUT SYSTEM

```
print("\n==== DRUG PREDICTION SYSTEM ===\n")
# Input Values
age = int(input("Enter Age: "))
sex = input("Enter Sex (M/F): ").strip().upper()
bp = input("Enter Blood Pressure (LOW/NORMAL/HIGH): ").strip().upper()
cholesterol = input("Enter Cholesterol Level (LOW/HIGH): "
").strip().upper()
na_to_k = float(input("Enter Sodium to Potassium Ratio: "))

# Get Prediction
predicted_drug = predict_drug(age, sex, bp, cholesterol, na_to_k)

# Display Output
print("====")
print(f"Predicted Drug: {predicted_drug}")
print("====\n")

==== DRUG PREDICTION SYSTEM ===

Enter Age: 30
Enter Sex (M/F): M
Enter Blood Pressure (LOW/NORMAL/HIGH): Low
Enter Cholesterol Level (LOW/HIGH): high
Enter Sodium to Potassium Ratio: 23.545

=====
Predicted Drug: drugY
=====
```